

## Load the dataset

In [114...]:

```
import pandas as pd
```

```
df = pd.read_csv("car_details.csv")
df.head()
```

Out[114...]:

	name	year	selling_price	km_driven	fuel	seller_type	transmission	owner
0	Maruti 800 AC	2007	60000	70000	Petrol	Individual	Manual	First Owner
1	Maruti Wagon R LXI Minor	2007	135000	50000	Petrol	Individual	Manual	First Owner
2	Hyundai Verna 1.6 SX	2012	600000	100000	Diesel	Individual	Manual	First Owner
3	Datsun RediGO T Option	2017	250000	46000	Petrol	Individual	Manual	First Owner
4	Honda Amaze VX i-DTEC	2014	450000	141000	Diesel	Individual	Manual	Second Owner

## Shape of the dataset

In [115...]:

```
df.shape
```

Out[115...]:

```
(4340, 8)
```

## Number of duplicates

In [116...]:

```
print(df.duplicated().sum())
```

763

## Drop the duplicates

In [117...]:

```
df = df.drop_duplicates()
df.shape
```

Out[117...]:

```
(3577, 8)
```

## Number of null values

In [118]: df.isnull().sum()

```
Out[118]: name      0
          year      0
          selling_price 0
          km_driven    0
          fuel        0
          seller_type   0
          transmission 0
          owner       0
          dtype: int64
```

## Data types

In [119]: df.dtypes

```
Out[119]: name      str
          year     int64
          selling_price  int64
          km_driven    int64
          fuel        str
          seller_type   str
          transmission str
          owner       str
          dtype: object
```

## Handling "name" column

In [120]: df["name"].value\_counts()

name	count
Maruti Swift Dzire VDI	54
Maruti Alto 800 LXI	48
Maruti Alto LXi	42
Maruti Alto LX	30
Hyundai EON Era Plus	28
..	
Maruti Swift LDI	1
Tata Nano XM	1
Mahindra Verito 1.5 D6 BSIII	1
Toyota Innova 2.5 VX (Diesel) 8 Seater BS IV	1
Hyundai i20 Magna 1.4 CRDi	1
Name: count, Length: 1491, dtype: int64	

In [121]: df["name"].nunique()

```
Out[121]: 1491
```

## Handling "fuel", "seller\_type", "transmission", "brand" columns

```
In [122... df["brand"] = df["name"].str.split(" ").str[0]
         df.drop("name", axis=1, inplace=True)
         df["brand"].value_counts()
```

```
Out[122... brand
          Maruti      1072
          Hyundai     637
          Mahindra    328
          Tata        308
          Ford        220
          Honda       216
          Toyota      170
          Chevrolet   151
          Renault     110
          Volkswagen  93
          Nissan      52
          Skoda       49
          Fiat         32
          Audi         31
          Datsun      29
          BMW         25
          Mercedes-Benz 21
          Jaguar      5
          Mitsubishi  5
          Land         5
          Volvo        4
          Jeep         3
          Ambassador   3
          MG           2
          OpelCorsa   2
          Daewoo       1
          Force        1
          Isuzu        1
          Kia          1
          Name: count, dtype: int64
```

```
In [123... df["brand"].nunique()
```

```
Out[123... 29
```

```
In [124... df["fuel"].value_counts()
```

```
Out[124... fuel
          Diesel      1800
          Petrol      1717
          CNG         37
          LPG         22
          Electric    1
          Name: count, dtype: int64
```

```
In [125... df["seller_type"].value_counts()
```

```
Out[125... seller_type
          Individual    2832
          Dealer        712
          Trustmark Dealer 33
          Name: count, dtype: int64
```

```
In [126... df["transmission"].value_counts()
```

```
Out[126... transmission
Manual      3265
Automatic    312
Name: count, dtype: int64
```

```
In [127... df = pd.get_dummies(df, columns=["fuel", "seller_type", "transmission", "
```

## Handling "owner" column

```
In [128... df["owner"].value_counts()
```

```
Out[128... owner
First Owner        2218
Second Owner       978
Third Owner        289
Fourth & Above Owner   75
Test Drive Car     17
Name: count, dtype: int64
```

```
In [129... owner_mapping = {
    "Test Drive Car": 0,
    "First Owner": 1,
    "Second Owner": 2,
    "Third Owner": 3,
    "Fourth & Above Owner": 4
}

df["owner"] = df["owner"].map(owner_mapping)
```

```
In [130... df.head()
```

```
Out[130...   year  selling_price  km_driven  owner  fuel_Diesel  fuel_Electric  fuel_LPG  fuel_Pel
0  2007        60000     70000      1          0            0            0            0
1  2007       135000     50000      1          0            0            0            0
2  2012       600000    100000      1          1            0            0            0
3  2017       250000     46000      1          0            0            0            0
4  2014       450000    141000      2          1            0            0            0
```

5 rows × 39 columns



## Split the dataset

```
In [131... from sklearn.model_selection import train_test_split
```

```
X = df.drop("selling_price", axis=1)
y = df["selling_price"]
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

print("Training data shape:", X_train.shape)
print("Testing data shape:", X_test.shape)
```

Training data shape: (2861, 38)  
 Testing data shape: (716, 38)

## Linear Regression Model

In [132...]

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, r2_score

linear_model = LinearRegression()
linear_model.fit(X_train, y_train)

linear_y_pred = linear_model.predict(X_test)

linear_mae = mean_absolute_error(y_test, linear_y_pred)
linear_r2 = r2_score(y_test, linear_y_pred)

print("Linear Regression\n")
print(f"MAE: {linear_mae:.4f}\n")
print(f"R2 Score: {linear_r2:.4f}")
```

Linear Regression

MAE: 180,314.2719

R2 Score: 0.5373

## Decision Tree Regressor Model

In [133...]

```
from sklearn.tree import DecisionTreeRegressor

tree_model = DecisionTreeRegressor(random_state=42)
tree_model.fit(X_train, y_train)

tree_y_pred = tree_model.predict(X_test)

tree_mae = mean_absolute_error(y_test, tree_y_pred)
tree_r2 = r2_score(y_test, tree_y_pred)

print("DecisionTreeRegressor\n")
print(f"MAE: {tree_mae:.4f}\n")
print(f"R2 Score: {tree_r2:.4f}")
```

DecisionTreeRegressor

MAE: 203,422.6013

R2 Score: 0.2580

## Random Forest Regressor Model

```
In [134...]: from sklearn.ensemble import RandomForestRegressor

forest_model = RandomForestRegressor(random_state=42)
forest_model.fit(X_train, y_train)

forest_y_pred = forest_model.predict(X_test)

forest_mae = mean_absolute_error(y_test, forest_y_pred)
forest_r2 = r2_score(y_test, forest_y_pred)

print("DecisionTreeRegressor\n")
print(f"MAE: {forest_mae:.4f}\n")
print(f"R2 Score: {forest_r2:.4f}")
```

DecisionTreeRegressor

MAE: 159,718.9318

R2 Score: 0.5842

## Gradient Boosting Regressor Model

```
In [135...]: from sklearn.ensemble import GradientBoostingRegressor

gradient_model = GradientBoostingRegressor(random_state=42)
gradient_model.fit(X_train, y_train)

gradient_y_pred = gradient_model.predict(X_test)

gradient_mae = mean_absolute_error(y_test, gradient_y_pred)
gradient_r2 = r2_score(y_test, gradient_y_pred)

print("DecisionTreeRegressor\n")
print(f"MAE: {gradient_mae:.4f}\n")
print(f"R2 Score: {gradient_r2:.4f}")
```

DecisionTreeRegressor

MAE: 157,631.8637

R2 Score: 0.5608

## Hyper Parameter Tuning the Random Forest Regressor

```
In [138...]: from sklearn.model_selection import RandomizedSearchCV
import numpy as np

random_grid = {
    'n_estimators': [100, 200, 300, 500],
    'max_features': ['sqrt', 'log2', None],
    'max_depth': [10, 20, 30, 40, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}
```

```
grid_search = RandomizedSearchCV(estimator=forest_model, param_distributi
grid_search.fit(X_train, y_train)

best_forest_model = grid_search.best_estimator_
best_y_pred = best_forest_model.predict(X_test)

print("Hyper Tuned Random Forest Regressor Model\n")
print(f"Best Parameters: {grid_search.best_params_}\n")
print(f"Best MAE: {mean_absolute_error(y_test, best_y_pred):,.4f}\n")
print(f"Best R2 Score: {r2_score(y_test, best_y_pred):.4f}")
```

Fitting 3 folds for each of 20 candidates, totalling 60 fits  
Hyper Tuned Random Forest Regressor Model

Best Parameters: {'n\_estimators': 200, 'min\_samples\_split': 2, 'min\_sample\_s\_leaf': 1, 'max\_features': 'log2', 'max\_depth': None}

Best MAE: 154,199.1606

Best R2 Score: 0.5519

## Save the final Random Forest Classifier Model

In [139...]

```
import joblib

joblib.dump(forest_model, "car_model.pkl")

model_columns = list(X.columns)
joblib.dump(model_columns, "model_columns.pkl")

print("Model and columns successfully saved!")
```

Model and columns successfully saved!